

Stock Analysis with Donald Trump Tweets

Harsh, Deepanshu Tyagi

May 7, 2017

STEP 1: Extract company names from Donald Trump's Tweets

First, we will load all of the libraries necessary to conduct this analysis.

```
library(plyr);library(dplyr);library(readr);library(twitterR);library(streamR)
;library(ROAuth);library(RCurl);library(stringr);library(stringdist);library(
pbapply);library(readxl);library(PerformanceAnalytics);library(xts);library(g
data);library(lubridate);library(tidyr);library(tidytext);
```

Next, read in authentication credentials to allow us to connect to the Twitter API.

```
twitter_cred <- read.csv("C:/Users/mnest/Google
Drive/RWD/Twitter/twitter_credentials.csv")
my_oauth <-
setup_twitter_oauth(twitter_cred[1,2],twitter_cred[2,2],twitter_cred[3,2],twi
tter_cred[4,2])
```

Get Tweets from the official Donald Trump Twitter account (@realDonaldTrump)

```
trump <- getUser("@realDonaldTrump")
#Tweets from API are stored in a list where each element of the list is an
individual tweet
trump_tweets <- userTimeline(trump, n=2000)
#Convert list of Tweets into a DataFrame
trump_tweets_df <- twListToDF(trump_tweets)
```

Read in all of the publicly traded company names from NASDAQ and NYSE. Remove common words from company names to increase accuracy of string distance computations.

```
nasdaq <- read_csv("C:/Users/mnest/Google Drive/Stevens/FE-
582(DataScience)/group_project/nasdaq.csv")
nyse <- read_csv("C:/Users/mnest/Google Drive/Stevens/FE-
582(DataScience)/group_project/nyse.csv")
stopwords <- "\\b(Corporation|Inc|Inc.|Corp|Co|Group|Resources|Systems)\\b"
nasdaq$Name_clean <- str_trim(str_replace_all(nasdaq$Name, stopwords, ""))
nyse$Name_clean <- str_trim(str_replace_all(nyse$Name, stopwords, ""))
```

Define a function to read each Trump tweet, strip all of the words in the tweet and compare them to the list of publicly traded company names. If the string similarity is within a predetermined threshold, then we accept this as a company that Trump has mentioned in his tweet.

```
get_trump_tweet_companies <- function(tweet){
  tmp_split <- unlist(str_split(toupper(iconv(enc2utf8(tweet)))), " "))

  trump_nyse_companies <- sapply(tmp_split, function(x) max(1 -
stringdist(x, toupper(nyse$Name_clean), method = "jw", p=.2))>.98)
  trump_nyse_companies_list <-
sapply(names(trump_nyse_companies)[trump_nyse_companies], function(x)
nyse$Name[amatch(x, nyse$Name_clean, method = "jw", p=.2, maxDist = 2000)])

  trump_nasdaq_companies <- sapply(tmp_split, function(x) max(1 -
stringdist(x, toupper(nasdaq$Name_clean), method = "jw", p=.2))>.98)
  trump_nasdaq_companies_list <-
sapply(names(trump_nasdaq_companies)[trump_nasdaq_companies], function(x)
nasdaq$Name[amatch(x, nasdaq$Name_clean, method = "jw", p=.2, maxDist =
2000)])

  results <- na.omit(c(trump_nyse_companies_list,
trump_nasdaq_companies_list))
  names(results) <- NULL

  return(unlist(results)[1])
}
```

Create a list that will contain all of the companies that Trump has mentioned in his tweets.

```
company_mentions <- pblapply(trump_tweets_df$text, get_trump_tweet_companies)
#Append this information to our original Trump Twitter DataFrame
trump_tweets_df$company_mentions <- company_mentions
#Append a logical column indicating if a given tweet contains a company name.
trump_tweets_df$is_populated <- lengths(trump_tweets_df$company_mentions) >=
1
```

STEP 2: Analyze the volatility of each stock Trump has mentioned in his tweets.

After the data was acquired from the Bloomberg terminal, we created an excel spreadsheet where each sheet contains information from the time when Trump mentioned a company. Then, we defined a function to read each sheet and plot the time series of respective Percent Return and Volume.

```
get_volatility_analysis <- function(trump_tweet_mention){
  trump_min_df <- read_excel("C:/Users/neste/Google Drive/Stevens/FE-
582(DataScience)/group_project/trump_minute_data.xlsx", sheet =
```

```

trump_tweet_mention, skip = 2)
  day_boundary=grep("9:30",trump_min_df$Date)

  day1=trump_min_df[1:day_boundary[2]-1, c("Date","LAST_PRICE", "VOLUME")]

  #Volume outlier detection, removes the first record, and all records
  following 3:59pm
  remove_me <- -grep(x = day1$Date, pattern = "(9[:]:30|(
16[:])|(15[:]:59))")

  day1_xts <- xts(day1$LAST_PRICE, order.by = day1$Date)

  day1_returns <- CalculateReturns(day1_xts)

  day1_return_pct <- day1_returns * 100

  day1_volume <- xts(day1$VOLUME[remove_me], order.by =
day1$Date[remove_me])

  trump_tweet <- unlist(read_excel("C:/Users/neste/Google
Drive/Stevens/FE-582(DataScience)/group_project/trump_minute_data.xlsx",sheet
= trump_tweet_mention)[1,8]);names(trump_tweet) <- NULL
  trump_tweet <- gsub(trump_tweet, pattern = "(AM|PM)", replacement = "")

  par(mfrow = c(1,2))
  plot(day1_return_pct, main = sprintf("Daily Return for %s",
trump_tweet_mention), ylab = "Percent Return")
  plot(day1_volume, main = sprintf("Volume Volatility for %s",
trump_tweet_mention), ylab = "Volume")
}

```

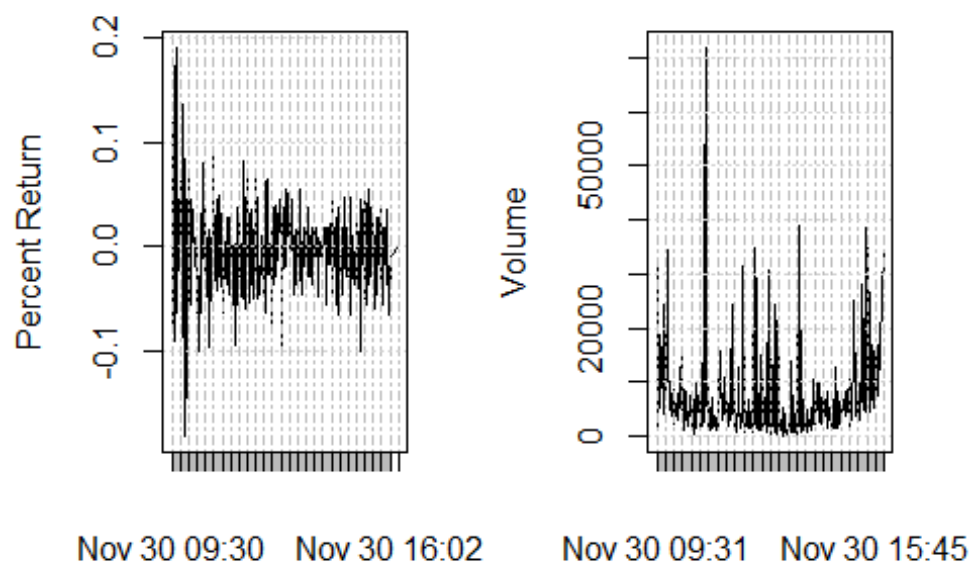
Next, we establish a character vector containint the tickers used in each sheet of our data, then run the function to generate the volatility plots.

```


#Character vector of tickers Trump mentioned in his tweets
trump_stock_mentions <- c("UTX1", "UTX2", "UTX3", "UTX4", "UTX5", "RXN",
"BA1", "BA2", "F1", "F2", "F3", "F4", "F5", "TWTR", "GOOG", "FB")
#For loop to return plots of each sheet
for(x in trump_stock_mentions){
  get_volatility_analysis(x)
}

```


Daily Return for UTX1 Volume Volatility for UTX1




Daily Return for UTX2 Volume Volatility for UTX2



Daily Return for UTX3 Volume Volatility for UTX3




Nov 30 09:30 Nov 30 16:02

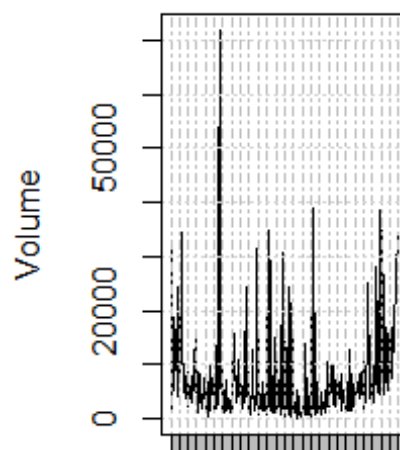


Nov 30 09:31 Nov 30 15:45

Daily Return for UTX4 Volume Volatility for UTX4




Nov 30 09:30 Nov 30 16:02




Nov 30 09:31 Nov 30 15:45

Daily Return for UTX5 Volume Volatility for UTX5




Nov 25 09:30 Nov 25 16:15




Nov 25 09:31 Nov 25 13:40

Daily Return for RXN Volume Volatility for RXN




Dec 05 09:30 Dec 05 16:02




Dec 05 09:31 Dec 05 15:45

Daily Return for BA1



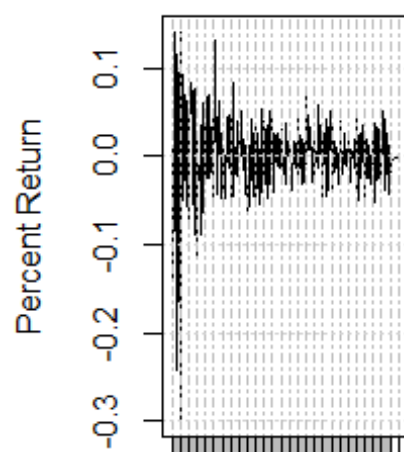
Feb 17 09:30 Feb 17 16:00

Volume Volatility for BA



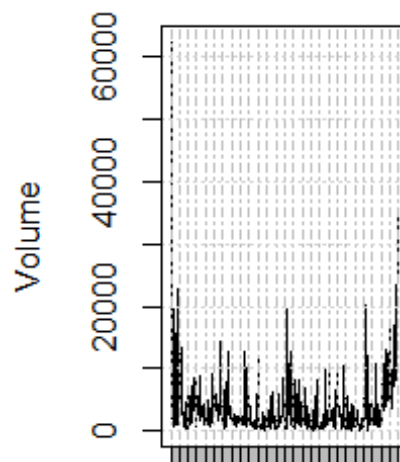
Feb 17 09:31 Feb 17 15:45

Daily Return for BA2




Dec 23 09:30 Dec 23 16:02

Volume Volatility for BA




Dec 23 09:31 Dec 23 15:45

Daily Return for F1



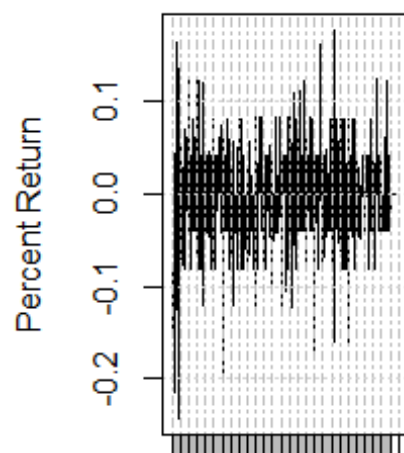
Jan 25 09:30 Jan 25 16:00

Volume Volatility for F1




Jan 25 09:31 Jan 25 15:45

Daily Return for F2




Jan 18 09:30 Jan 18 16:01

Volume Volatility for F2




Jan 18 09:31 Jan 18 15:45

Daily Return for F3




Jan 09 09:30 Jan 09 16:03

Volume Volatility for F3



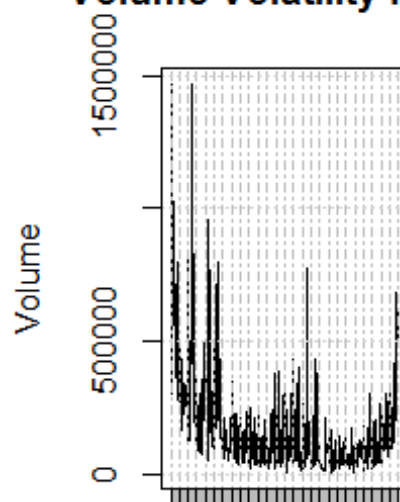
Jan 09 09:31 Jan 09 15:45

Daily Return for F4




Jan 04 09:30 Jan 04 16:00

Volume Volatility for F4




Jan 04 09:31 Jan 04 15:45

Daily Return for F5




Volume Volatility for F5




Nov 17 09:30 Nov 17 16:00

Nov 17 09:31 Nov 17 15:45

Daily Return for TWTR




Volume Volatility for TWTR




Oct 31 09:30 Oct 31 15:45

Oct 31 09:31 Oct 31 15:30

Daily Return for GOOG Volume Volatility for GOOG



Daily Return for FB Volume Volatility for FB



STEP 3: Utilize the sentiment from Donald Trump's tweets to ultimately create an ETF.

First, we calculated the word frequencies within Donald Trump's tweets

```
trump_tweets_df <- read_csv('C:/Users/neste/Google Drive/Stevens/FE-582(DataScience)/group_project/trump_tweets.csv')
data("stop_words")
trump_tweets_tidy <- trump_tweets_df %>%
  mutate(text = toupper(iconv(text))) %>%
  unnest_tokens(output = word, text) %>%
  anti_join(stop_words) %>%
  count(word, sort=T) %>%
  ungroup()

head(trump_tweets_tidy, 20)
```

```
## # A tibble: 20 × 2
##       word      n
##   <chr> <int>
## 1  https    379
## 2   t.co    379
## 3   amp     89
## 4  people    64
## 5   trump    52
## 6  america    46
## 7 draintheswamp 46
## 8   time    46
## 9   join    44
## 10   u.s     44
## 11  country    42
## 12  hillary    42
## 13  election    41
## 14   media    41
## 15  clinton    40
## 16   jobs     39
## 17    â       38
## 18   news     38
## 19  president    36
## 20   fake     33
```

Next, we needed to establish a positive/negative dictionary to evaluate the sentiment of each tweet. To do this, we incorporated the most frequently used words in each of Trump's tweets which we personally marked as positive or negative, as well as the Stanford NLP dictionary of positive and negative words.

```
#Read in Team 7's modified sentiment dictionary based on the word frequency in Trump Tweets
tmp <- read_csv('C:/Users/neste/Google Drive/Stevens/FE-582(DataScience)/group_project/trump_freq_words.csv')
```

```

colnames(tmp) <- c("word", "sentiment")
tmp <- tmp[!is.na(tmp$sentiment), ]
nrow(tmp)

## [1] 57

tmp_list <- lapply(split(tmp, tmp$sentiment), function(x) x$word)
tmp_negative <- tmp_list$`0`
length(tmp_negative)

## [1] 32

tmp_positive <- tmp_list$`1`
length(tmp_positive)

## [1] 25

positive_words <- toupper(unlist(read.delim(file = 'C:/Users/neste/Google
Drive/RWD/NLP/positive-words.txt', stringsAsFactors =
F)));names(positive_words) <-NULL
positive_words <- positive_words[-c(1:32)]
print(sprintf("The number of positive words in this vector is:
%d",length(positive_words)))

## [1] "The number of positive words in this vector is: 2006"

positive_words <- unique(append(positive_words, toupper(tmp_positive)))
print(sprintf("The number of positive words in this vector is NOW:
%d",length(positive_words)))

## [1] "The number of positive words in this vector is NOW: 2024"

negative_words <- toupper(unlist(read.delim(file = 'C:/Users/neste/Google
Drive/RWD/NLP/negative-words.txt', stringsAsFactors =
F)));names(negative_words) <-NULL
negative_words <- negative_words[-c(1:32)]
print(sprintf("The number of negative words in this vector is:
%d",length(negative_words)))

## [1] "The number of negative words in this vector is: 4783"

negative_words <- unique(append(negative_words, toupper(tmp_negative)))
print(sprintf("The number of negative words in this vector is NOW:
%d",length(negative_words)))

## [1] "The number of negative words in this vector is NOW: 4805"

```

Read in the Donald Trump tweets DataFrame, clean-up the created date/time column and convert to a date variable.

```
trump_tweets_df <- read_csv("C:/Users/neste/Google
Drive/RWD/Twitter/trump/trump_tweets.csv") %>%
  separate(created, c("Date", "Time"), sep = " ", remove = T) %>%
  mutate(Date = ymd(Date))
```

Define two functions. First function will compute the sentiment score for given text. Second function will return the number of capital letters for given text.

```
get_sentiment_score <- function(a){
  #split the string by spaces to extract each individual word from the
  given text
  a_split <- unlist(strsplit(toupper(iconv(a)), " "))

  #positive count will be equal to the length of the set of words
  contained within the positive dictionary and our given words.
  positive_count <- length(intersect(a_split, positive_words))
  #negative count will be equal to the length of the set of words
  contained within the negative dictionary and our given words.
  negative_count <- length(intersect(a_split, negative_words))

  return(positive_count / (positive_count + negative_count))
}

get_capital_letters <- function(a){
  a_split <- unlist(strsplit(iconv(a), NULL))
  return(sum(str_detect(a_split, "[A-Z]"), na.rm = T))
}
```

For each day that Donald Trump tweets, generate the average sentiment score of his tweets. If a score happens to be NA, we will impute it to be the average of all daily sentiment scores.

```
trump_sentiment_df <- trump_tweets_df %>%
  rowwise() %>%
  mutate(sentiment_score = get_sentiment_score(text)) %>%
  mutate(capital_score = get_capital_letters(text)) %>%
  ungroup() %>%
  group_by(Date) %>%
  summarise(daily_sentiment_score = mean(sentiment_score, na.rm = T),
    total_capital_letters = sum(capital_score, na.rm=T), retweets_favorites =
    sum(c(retweetCount, favoriteCount), na.rm=T))

#Impute NA scores to be the average daily sentiment score.
trump_sentiment_df$daily_sentiment_score[is.na(trump_sentiment_df$daily_senti
ment_score)] <- mean(trump_sentiment_df$daily_sentiment_score, na.rm=T)
#Show 15 records of results
head(trump_sentiment_df, 15)
```

```
## # A tibble: 15 × 4
##       Date daily_sentiment_score total_capital_letters
##       <date>                <dbl>                <int>
## 1  2016-10-17                0.3484848                116
## 2  2016-10-18                0.6388889                184
## 3  2016-10-19                0.3076923                313
## 4  2016-10-20                0.5183908                662
## 5  2016-10-21                0.4861111                187
## 6  2016-10-22                0.6346154                245
## 7  2016-10-23                0.0000000                 15
## 8  2016-10-25                0.4444444                165
## 9  2016-10-26                0.3333333                 68
## 10 2016-10-27                0.6428571                303
## 11 2016-10-28                0.4047619                121
## 12 2016-10-29                0.6666667                124
## 13 2016-10-30                0.6979167                128
## 14 2016-10-31                0.5833333                 57
## 15 2016-11-02                1.0000000                 23
## # ... with 1 more variables: retweets_favorites <dbl>
```

Read in the ticker list of nasdaq and nyse.

```
ticker_list_nyse <- read_csv("C:/Users/neste/Google Drive/Stevens/FE-
582(DataScience)/group_project/nyse.csv") %>%
  select(Symbol) %>%
  unlist
names(ticker_list_nyse) <- NULL

ticker_list_nasdaq <- read_csv("C:/Users/neste/Google Drive/Stevens/FE-
582(DataScience)/group_project/nasdaq.csv") %>%
  select(Symbol) %>%
  unlist
names(ticker_list_nasdaq) <- NULL
```

Define function to read ticker information from yahoo.

```
yahoo_read <- function(url){
  require(RCurl)
  if(url.exists(url)){
    dat <- read.table(url,header=TRUE,sep=",")
    df <- dat[,c(1,5)]
    df$Date <- as.Date(as.character(df$Date))
    return(df)
  }
}
```

Define a function to read the ticker information from yahoo, calculate the return, then run a multiple linear regression against the trump sentiment data.

```
trump_ticker_fit <- function(ticker){
  cat(ticker); cat(' ')
  ticker_url <- paste0(paste0('http://real-
chart.finance.yahoo.com/table.csv?s=', ticker, '&a=07&b=24&c=2010&d=12&e=22&f=2
016&g=d&ignore=.csv'))

  ticker_df <- try(yahoo_read(ticker_url), silent = T)
  if(any(class(ticker_df) == "try-error", nrow(ticker_df) < 20,
is.null(ticker_df))){
    return(NA)
  }
  ticker_df$returns <- CalculateReturns(ts(ticker_df$Close))

  trump_vs_ticker_df <- trump_sentiment_df %>%
    inner_join(select(ticker_df, Date, returns), by = "Date")
  cor(trump_vs_ticker_df[, -1])

  fit <- try(lm(returns ~., data = trump_vs_ticker_df[, -1]), silent = T)
  if(class(fit) == "try-error"){
    return(NA)
  }
  rsq <- summary(fit)$r.squared
  return(rsq)
}
```

Run the model for each ticker in NYSE and NASDAQ, then upload all of this information into a single DataFrame.

```
stocks_to_invest <- pbsapply(c(ticker_list_nyse, ticker_list_nasdaq),
trump_ticker_fit)

stocks_to_invest_df <- data.frame(ticker = names(stocks_to_invest), rsq =
stocks_to_invest, stringsAsFactors = F)
```